

HPC Workshop Series: *Compilation*

Nick Featherstone
Research Computing



Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

Objects



Outline

- Makefile basics
- Creating Shared & Static Libraries
- Linking Shared & Static Libraries
- Environment considerations
- Linking order



Before we begin

- `ssh username@tutorial-login.rc.colorado.edu`
- `ssh scompile`
- `ml intel`



Nano survival guide

- Ctrl+o save (need to confirm filename)
- Ctrl+x exit
- Ctrl+k cut
- Ctrl+u paste



GNU Make

- Helps manage compilation of source code.
- Examine the files in the makefiles subdirectory
- Compile:
 - `icc hello.c hellofunc.c -I. -o hello.exe`
- Drawbacks?
 - Have to remember compilation syntax
 - Possibly machine-specific (no Intel on my laptop?)



Simple Makefile (1)

- Create a file named Makefile in the makefiles directory
- Add these lines (use tab, not 4 spaces):

```
hello.exe: hello.c hellofunc.c
```

```
TAB icc -I . hello.c hellofunc.c -o hello.exe
```

- This is a **rule** for making the **target file** hello.exe
- Type ***make***
- Checks target and dependencies for:
 - existence
 - modifications (based on timestamp)



Simple Makefile (1)

- Make executes every command indented below the target/dependency line.
- Can use any Linux commands

```
hello.exe: hello.c hellofunc.c
```

```
TAB echo "Compiling hello.exe"
```

```
TAB icc -I . hello.c hellofunc.c -o hello.exe
```

- Make displays the command executed on the screen
- Suppress by prepending @ symbol: echo → @echo



Simple Makefile (1)

- Standard to include a target named clean:

```
hello.exe: hello.c hellofunc.c
```

```
TAB @echo "Compiling hello.exe"
```

```
TAB icc -I . hello.c hellofunc.c -o hello.exe
```

```
clean:
```

```
TAB rm -f *.o
```

```
TAB rm -f hello.exe
```



Simple Makefile (1)

- Complication. Run these commands:
 - touch clean
 - make clean
- A file named clean will conflict with our target named clean.
- Solution – add .PHONY target:

```
.PHONY: clean
```

- Understood that targets in .PHONY list are not files



Makefile (1): A Good Start

```
.PHONY: clean
```

```
hello.exe: hello.c hellofunc.c
```

```
TAB @echo "Compiling hello.exe"
```

```
TAB icc -I . hello.c hellofunc.c -o hello.exe
```

```
clean:
```

```
TAB rm -f *.o
```

```
TAB rm -f hello.exe
```



Variables in Make

- Compilers and optimization flags may differ between machines
- Use variables to avoid rewriting compilation commands for each new machine
- Best to do at top of Makefile

Define variables via **=**

```
CC = icc
```

Evaluate variables via **\$()**

```
echo $(CC)
```



Simple Makefile (2)

- Modify your makefile to include these changes
- (keep clean and .PHONY as they were)

```
CC = icc
INCLUDE_FLAGS = -I .
OPT_FLAGS = -O2
CFLAGS = $(INCLUDE_FLAGS) $(OPT_FLAGS)

hello.exe: hello.c hellofunc.c
    @echo "Compiling hello.exe"
    $(CC) $(CFLAGS) . hello.c hellofunc.c -o hello.exe
```



Special Variables

- These variables are useful within rules/recipes
- The variable $\$@$
 - Refers to the rule target (e.g., hello.exe)
- The variable $\$<$
 - Refers to the 1st dependency (e.g., hello.c)
- The variables $\$^$
 - Refers to all dependencies (e.g., hello.c hellofunc.c)
- Exercise:
 - Rewrite the rule for hello.exe using $\$@$ and $\$^$



My Solution

I defined a new variable...

```
...  
CFLAGS = $(OPT_FLAGS) $(INCLUDE_FLAGS)  
PROG = hello.exe
```

Now hello.exe appears in only 1 place.

```
$(PROG): hello.c hellofunc.c  
    @echo "Compiling " $(PROG) "  
    $(CC) $(CFLAGS) $^ -o $@  
  
clean:  
    rm -f *.o  
    rm -f $(PROG)
```



Simple Makefile (3)

- We often have a number of object files that needed to build the program.
- Make this small modification:

```
$(PROG): hello.c hellofunc.c
```

```
$(PROG): hello.o hellofunc.o
```



- By default, Make uses CFLAGS and CC to build the .o files
- We can control this behavior ourselves



Wildcards

- The % symbol acts a wildcard.
- Add the following target/rule:

```
%o: %.c  
    @echo "Compiling "$@ " using " $<  
    $(CC) $(CFLAGS) -c $< -o $@
```

- This tells Make:
 - Use this rule if a file ending in .o is needed
 - Before file.o can be created, file.c must exist
 - Build file.o via: \$(CC) \$(CFLAGS) -c file.c -o file.o



Portability

- If we move to a new machine, we may need to modify:
 - CC, OPT_FLAGS, INCLUDE_FLAGS
- To do so, we must edit the Makefile
- OK enough for small projects; really bad otherwise
- Good practice:
 - Create machine-specific definitions file
 - Include in Makefile



Portability

- Create a file named machine.def with these three lines

```
CC = icc  
OPT_FLAGS = -O2  
INCLUDE_FLAGS = -I .
```

- In your Makefile, make this replacement:

```
CC = icc  
OPT_FLAGS = -O2  
INCLUDE_FLAGS = -I .
```



```
include machine.def
```

- Makefile works on any machine now
- only machine.def is modified



II. Library Creation & Linking

Outline

- Compilation vs. Linking
- Shared vs. Static Libraries
- Creating static and shared libraries
- LD_LIBRARY_PATH
- ldd and nm
- RPATH
- Link order



Compiling vs. Linking

- Compilation generates object files (machine code) from source code
- Linking bundles object files together to generate a complete executable
- Often carried out via single command for small projects.
- Multiple stages for complex projects



Compilation/Linking Examples

Single-Step

```
icc -l . hello.c hellofunc.c -o hello.exe
```

Separate Steps

```
icc -l . -c hellofunc.c      -o hellofunc.o  
icc -l . -c hello.c          -o hello.o
```

} compilation

```
icc hello.o hellofunc.o -o hello.exe
```

} linking



Peering Behind the Curtain

- Quite a bit going on in the background
- Quick exercise:
 - In makefiles directory, edit machine.def
 - Change `-O2` to `-O2 -v`
 - run `make clean / make`
- “-v” denotes “verbose”



Libraries

- Common practice to bundle related functionality into single library
- Can be used in multiple programs
- Two flavors
 - Shared (dynamically linked at runtime)
 - Static (bundled into executable at compile time)



Shared Libraries

- .so extension (e.g., libc.so)
- Advantages
 - Easy to manipulate via environment
 - HPC centers often provide optimized versions of commonly-used libraries.
 - Smaller executable size
 - RAM efficient (multiple processes can access same shared library)
- Disadvantages:
 - occasional version conflicts
 - User must be “environment-aware”
- Advice:
 - Use shared libraries when incorporating 3rd party software into your project.



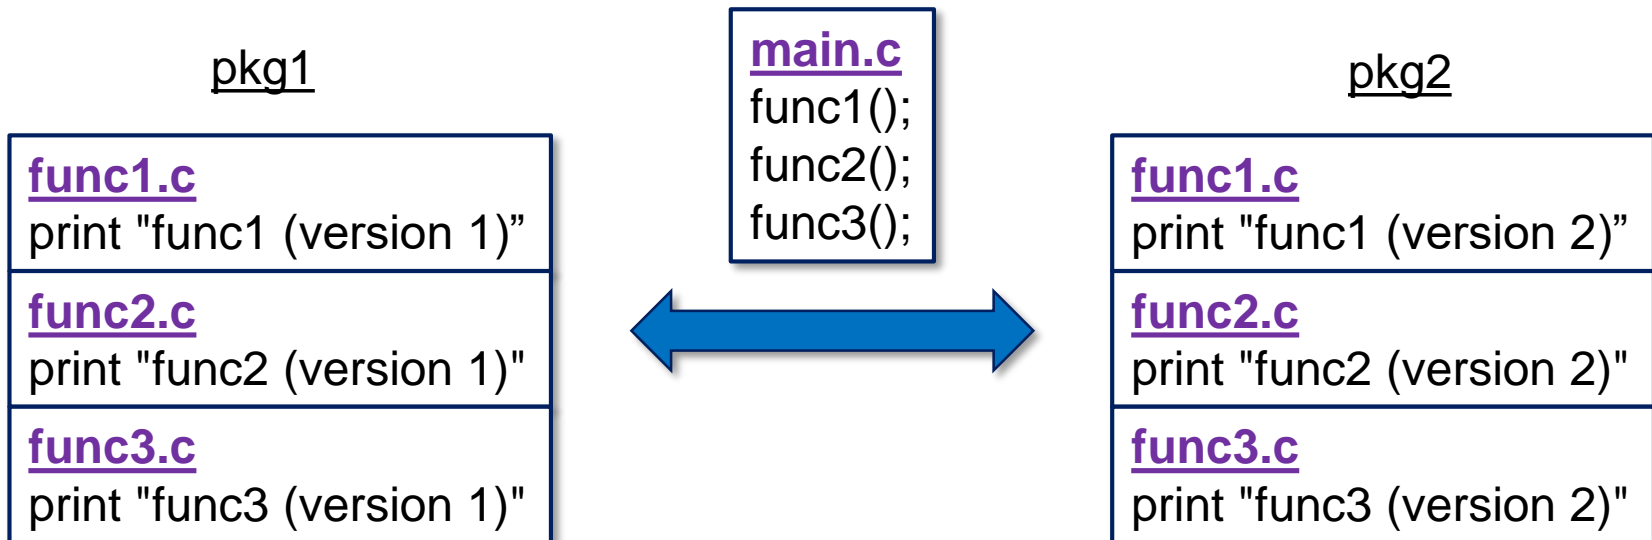
Static Libraries

- .a extension (e.g., liblapack.a)
- Advantages
 - More difficult to introduce version conflicts
 - Conflicts discovered at compiled time
- Disadvantages:
 - Larger executable size
 - Potential for redundant code in memory
 - More difficult to incorporate updates
- Advice:
 - Where it seems natural, use static libraries to bundle portions of your own project.
 - Use static libraries for 3rd party packages that tend to change significantly between versions (e.g., Boost)



Hands-On Exercise

- We are going to work with a small project that uses statically- and dynamically-linked libraries that we create.
- We will examine how to switch between shared libraries in pkg1 and pkg2 at run-time



Our pkg1 Makefile

- Change to the directory “libraries” this is our project directory.
- Open pkg1/src/Makefile
- Let’s have a look.
- Presently, instructions for creating the shared and static library are missing...



Creating a Static Library

- Step1: Compile object files needed by library

```
icc -I . -O1 -c func1.c      -o func1.o  
icc -I . -O1 -c func2.c      -o func2.o  
icc -I . -O1 -c func3.c      -o func3.o
```

*type this at
command line*

- Step2: Bundle your .o files using the archiver

```
ar rc libstatic.a func1.o func2.o func3.o
```

- GNU archiver options:
 - r – replace functions if already found to exist in library
 - c – create static library if it does not already exist



On Your Own

- Modify the (now empty) rule for `$(LIBSTATIC)` in your Makefile so that it:
 - Compiles the `func.o` files
 - Creates an archive from the `func.o` files
 - Use the variables defined in the Makefile and in `machine.def`
- Test it:
 - `make libstatic.a`



Creating a Static Library

- Step1: Compile object files with -fPIC

```
icc -I . -O1 -fPIC -c func1.c      -o func1.o  
icc -I . -O1 -fPIC -c func2.c      -o func2.o  
icc -I . -O1 -fPIC -c func3.c      -o func3.o
```

*type this at
command line*

- Step2: Use icc with -shared flag to create .so file

```
icc -shared func1.o func2.o func3.o -o libshared.so
```

- Shared library options:
 - -fPIC – generate Position Independent Code
 - -shared – create a shared library



On Your Own

- Modify the (now empty) rule for `$(LIBSHARED)` in your Makefile so that it:
 - Compiles the `func.o` files
 - Creates a shared library from the `func.o` files
 - Use the variables defined in the Makefile and in `machine.def`
- Test it:
 - `make libshared.so`
 - `make clean`
 - `make all` (will build both libs and “install”)



Creating package 2

- Change back to the “libraries” directory
- `cp -r pkg1 pkg2` (create pkg2 directory)
- Change to `pkg2/src`
- Modify each `funcX.c` to say “version 2”
- `make all`
- Change back to the “libraries” directory



Linking our Libraries

- Static and shared libraries are linked with different syntax at compile time
- Static libraries – provide full path to library:

```
icc -O2 -I . -I pkg1/include main.c pkg1/lib/libstatic.a -o test.static
```

- Shared libraries:
 - Indicate directory with `-L` flag
 - Indicate library using “l” shorthand (libshared.so = `-lshared`)

```
icc -O2 -I . -I pkg1/include main.c -Lpkg1/lib -lshared -o test.dynamic
```

- Try this at the command line



On Your Own

- Let's examine the Makefile in “libraries”
- Edit the LIBSTATIC and LIBSHARED variables so that test.dynamic and test.static build correctly.
- Run make clean & make to build
- Try it out:
 - ./test.static
 - ./test.dynamic
- What happens?



LD_LIBRARY_PATH

- Test.dynamic failed because the system did not know where to look for libshared
- Using dynamic libraries is a two-step process
 - Compile time: compiler pointed to library via -L
 - Run time: system checks system defaults and LD_LIBRARY_PATH directories for shared libraries
- LD_LIBRARY_PATH
 - Colon-separated list of directories
 - System checks first directory, then second, and so on when loading shared libraries at runtime.



LD_LIBRARY_PATH

- Example:
 - LD_LIBRARY_PATH = /lib:/home/mylib
 - System checks
 1. /lib first
 2. /home/mylib second
- Example:
 - LD_LIBRARY_PATH = /home/mylib:/usr/lib64
 - System checks
 1. /home/mylib first
 2. /lib second



Running Test.dynamic

- Save original LD_LIBRARY_PATH:

```
export LD_SAVE=$LD_LIBRARY_PATH
```

- Try this:

```
export LD_LIBRARY_PATH=pkg1/lib:$LD_SAVE  
./test.dynamic
```

- And this:

```
export LD_LIBRARY_PATH=pkg2/lib:$LD_SAVE  
./test.dynamic
```



Readelf: What is Needed?

- ELF = Executable Linkable Format
- Standard format for Linux executables and libraries
- Readelf: utility for parsing ELF files
- Tells us which dynamic libraries are needed
- readelf -h filename *view header*
- readelf -d filename *view dynamic section*
- Try it with test.dynamic and test.static



Readelf Output

readelf -d test.static

Dynamic section at offset 0x35c0 contains 27 entries:

Tag	Type	Name/Value
0x0000000000000001	(NEEDED)	Shared library: [libm.so.6]
0x0000000000000001	(NEEDED)	Shared library: [libgcc_s.so.1]
0x0000000000000001	(NEEDED)	Shared library: [libc.so.6]
0x0000000000000001	(NEEDED)	Shared library: [libdl.so.2]

readelf -d test.dynamic

Dynamic section at offset 0x35c0 contains 28 entries:

Tag	Type	Name/Value
0x0000000000000001	(NEEDED)	Shared library: [libshared.so]
0x0000000000000001	(NEEDED)	Shared library: [libm.so.6]
0x0000000000000001	(NEEDED)	Shared library: [libgcc_s.so.1]
0x0000000000000001	(NEEDED)	Shared library: [libc.so.6]
0x0000000000000001	(NEEDED)	Shared library: [libdl.so.2]



LDD: Where am I Pointed?

- The ldd utility indicates which shared library will be used
- Depends on system defaults and current LD_LIBRARY_PATH
- Try it out:
 - ldd test.static
 - ldd test.dynamic



Idd Output

Idd test.static

```
linux-vdso.so.1 => (0x00007fff584c1000)
libm.so.6 => /lib/x86_64-linux-gnu/libm.so.6 (0x00007fe0d3e40000)
libgcc_s.so.1 => /lib/x86_64-linux-gnu/libgcc_s.so.1 (0x00007fe0d3c2a000)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007fe0d3860000)
libdl.so.2 => /lib/x86_64-linux-gnu/libdl.so.2 (0x00007fe0d365c000)
/lib64/ld-linux-x86-64.so.2 (0x00007fe0d4149000)
```

Idd test.dynamic

```
linux-vdso.so.1 => (0x00007ffd11ddd000)
libshared.so => pkg1/lib/libshared.so (0x00007f379221b000)
libm.so.6 => /lib/x86_64-linux-gnu/libm.so.6 (0x00007f3791f12000)
libgcc_s.so.1 => /lib/x86_64-linux-gnu/libgcc_s.so.1 (0x00007f3791cfc000)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f3791932000)
libdl.so.2 => /lib/x86_64-linux-gnu/libdl.so.2 (0x00007f379172e000)
... (various Intel libraries here)
/lib64/ld-linux-x86-64.so.2 (0x00007f379241c000)
```



Exercise

- Sometimes dynamic linking creates unexpected problems.
- Copy pkg1 to a new directory pkg3
- Modify the FUNCS variable to omit func3.o
- Rebuild the library via: make clean / make
- Point LD_LIBRARY_PATH at pkg3/lib and re-run test.dynamic



Checking Symbols with nm

- Can use the nm utility to check symbols contained within a file
- e.g., nm pkg1/lib/libshared.so
- Examine the pkg1 and pkg3 libraries
- pkg1 shows the func3 symbol
- pkg3 does not



RPATH: Hard-coded Links

- If desired, we can both:
 - link dynamically
 - ignore LD_LIBRARY_PATH
- To do so, specify the RPATH when compiling, e.g.

```
-Wl,-rpath,$(PWD)/pkg1/lib
```

- RPATH supercedes LD_LIBRARY_PATH
- Examine libraries/Makefile



Running test.rpath

- Save original LD_LIBRARY_PATH:

```
export LD_SAVE=$LD_LIBRARY_PATH
```

- Try this:

```
export LD_LIBRARY_PATH=pkg1/lib:$LD_SAVE  
./test.rpath
```

- And this:

```
export LD_LIBRARY_PATH=pkg2/lib:$LD_SAVE  
./test.rpath
```



The **ORIGIN** Variable

- The **ORIGIN** variable allows us to specify a relative **RPATH**
- This makes it easier to copy code and shared libraries around.
- Compiler flags (double \$ is correct):

```
-Wl,-rpath,"$$$$ORIGIN"/lib
```

- Run: `readelf -d test.origin`



ORIGIN: Try it Out

- We can now move our executable and library wherever we like
- Test.origin will look for .so file in lib subdirectory

```
mkdir $USER/origin
```

```
mkdir $USER/origin/lib
```

```
cp test.origin $USER/origin/.
```

```
cp {pkg1 OR pkg2}/lib/libshared $USER/origin/lib  
/$USER/origin/test.origin
```



Link Order

- Can link multiple libraries when compiling a program. E.g.,
`ifort main.f90 -o prog1 -lv1 -lv2 -lv3`
- If libraries contain unique subroutines, order does not matter.
- If two or more libraries contain same subroutine names, it matters.
- Why would you have the same routine name?
 - Example: LAPACK & IBM ESSL
 - ESSL & LAPACK both have DGEMM
 - Only LAPACK has dgetrf (band solve routine)
 - May want to use ESSL DGEMM along with LAPACK dgetrf



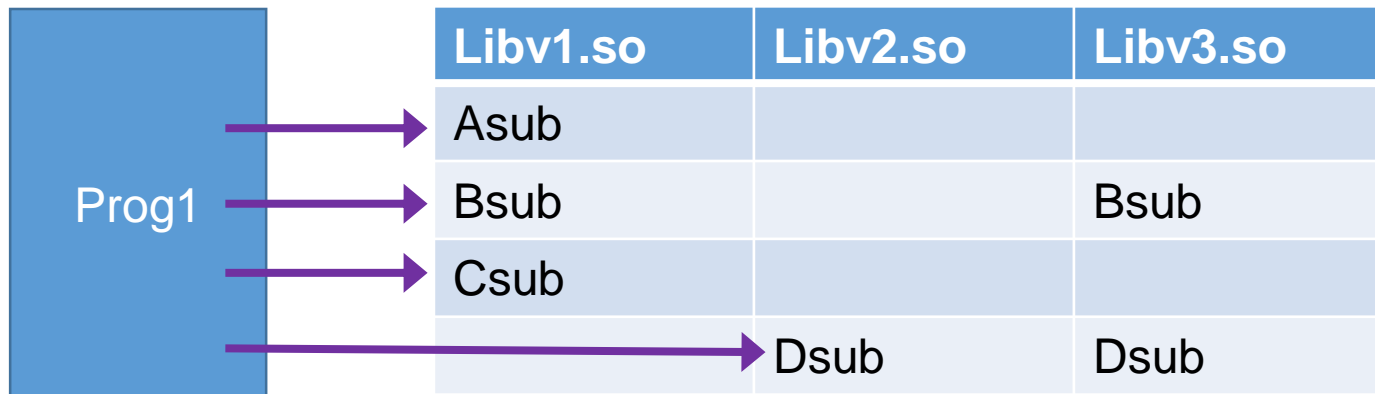
Link Order

- Let's examine the Makefile in the link_order directory
- Build the code: `make clean / make`
- Export `LD_LIBRARY_PATH=lib:$LD_LIBRARY_PATH`
- Run `prog1: ./prog1`



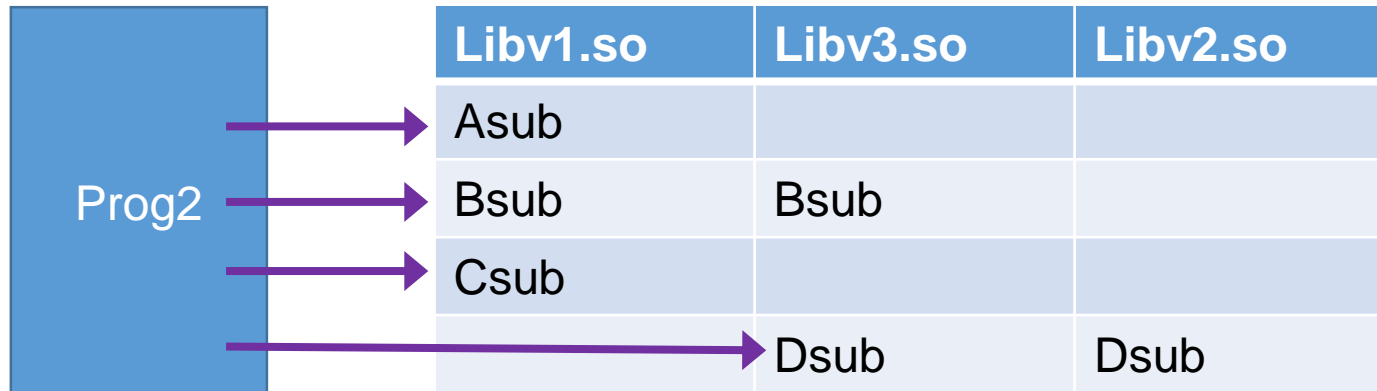
Link Order

- How does order of linking matter?
- Most linkers check from left to right
- Stops once symbol is found
- Good practice: put “root” dependency last
- Ex: `ifort main.f90 -o prog1 -lv1 -lv2 -lv3`



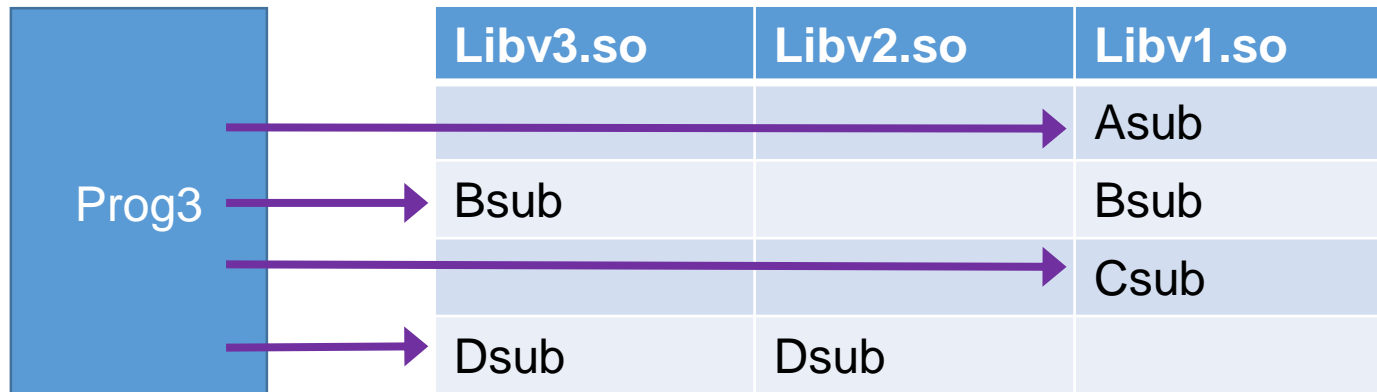
Link Order

- Run prog2...
 - `ifort main.f90 -o prog2 -lv1 -lv3 -lv2`



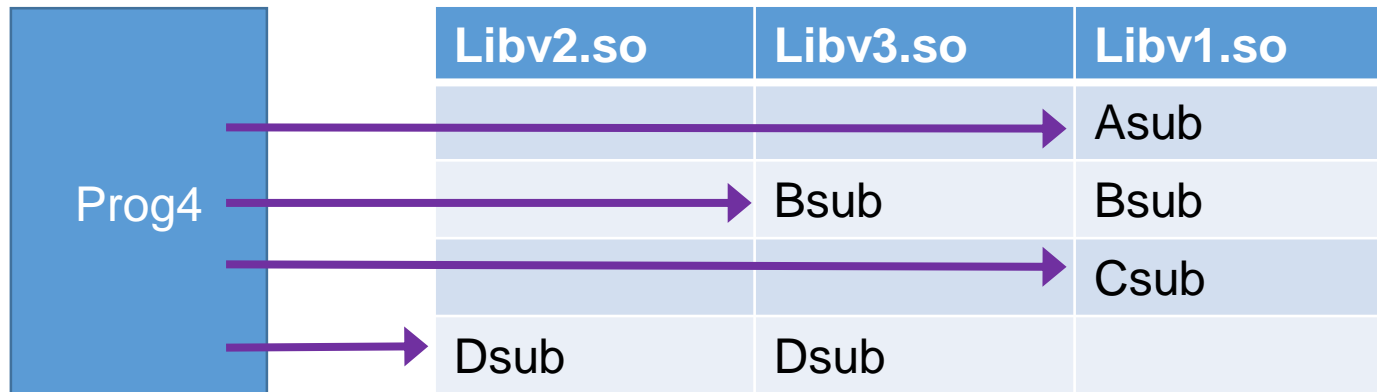
Link Order

- Run prog3...
 - `ifort main.f90 -o prog2 -lv1 -lv3 -lv2`



Link Order

- Run prog4...
 - `ifort main.f90 -o prog2 -lv2 -lv3 -lv1`



Test Yourself

- What will be the output of
 - `ifort main.f90 -o prog1 -lv2 -lv1 -lv3`
 - `ifort main.f90 -o prog1 -lv3 -lv1 -lv2`

Libv1.so	Libv2.so	Libv3.so
Asub		
Bsub		Bsub
Csub		
	Dsub	Dsub

